

# Big Data Engineering and Architecture

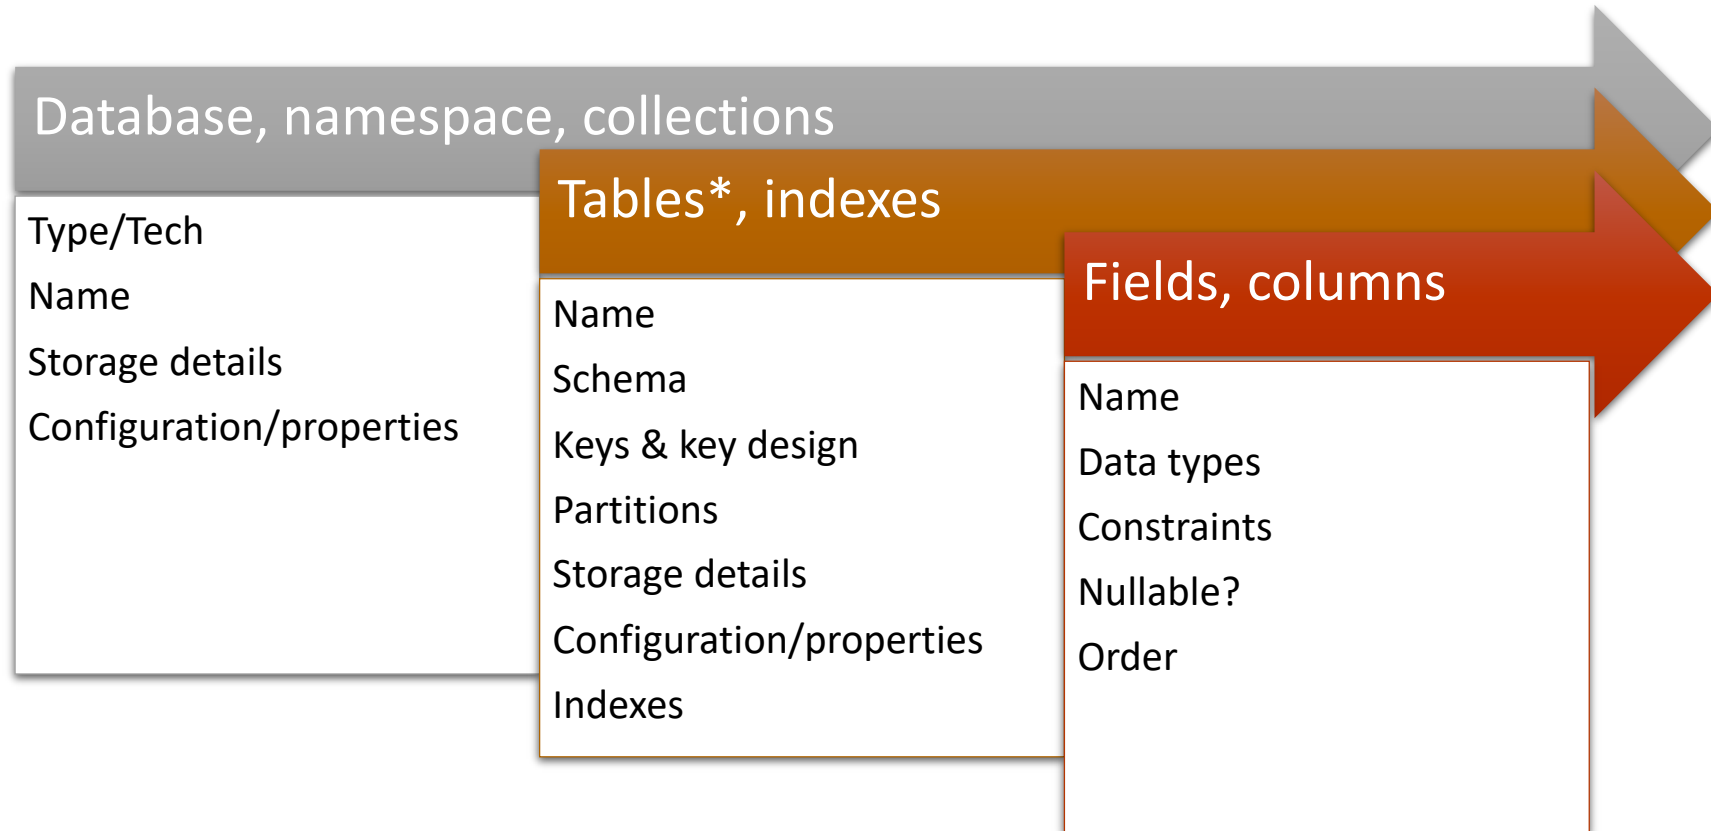
Topic 6: NoSQL Data Modeling (Key/Value and Document)

# What is Data Modeling?

- Process of creating a data model for data to be stored in a database
- Contains
  - Data objects
  - Associations between objects
  - Rules
- Types of data models in a design:
  - Conceptual
  - Logical
  - Physical (we will focus on this)

# Physical Data Model

Describes HOW the system will be implemented in the database



*\*NOTE for graphs: Tables would be nodes and arcs*

# Data Model Design

- Design the model(s) and document them
- Describe important characteristics
- Describe design considerations (why certain decisions were made)
- Use graphics wherever possible
- Use tables to describe attributes of each level
- Include examples and hints for users to anchor their understanding in instances

Key/Value Stores



# Overview

- Three essential features
  - Simplicity
  - Speed
  - Scalability
- Simplicity
  - Don't need all features of RDBMS (joins, multiple entity queries, schema)
- Speed
  - Simple data structure/access → Fast reads/writes
- Scalability
  - Scale out with minimal disruption

# Definitions

- Key
  - Reference to a value (like an address)
- Value
  - Data stored as referenced by a key (can be just about anything)
- Namespace
  - Collection of key-value pairs
  - Could be based on domain, database, or the entire key-value store
- Partition
  - Subset of a database
  - Store keys in different partitions based on their value

# Data Samples

Key	Value
cust_1234_name	"Peter Griffin"
cust_1234_address	"31 Spooner St. Quahog, RI"
cust_1234_nameAddress	{"name": "Peter Griffin", "address": "31 Spooner St. Quahog, RI"}
cust_1234_profileImage	
cust_2468_name	"Homer Simpson"
cust_2468_address	"742 Evergreen Terrace, Springfield, ???"
cust_2468_nameAddress	{"name": "Homer Simpson", "address": "742 Evergreen Terrace, Springfield, ???"}
cust_2468_profileImage	



# Data Structure Overview

- Data structure is an associative array in persistent storage
  - Dictionary, map, hash table
- Key-Value
  - Key is a unique identifier within the namespace
  - Value can be essentially anything (number, string, binary, etc.)
- Schemaless
  - Can hold multiple representations of the same data simultaneously
    - FirstName, LastName, FullName

# Keys

- Unique identifier within a namespace
- Can use Id value (like SSN), but usually use compound keys
  - Id.Attribute: 1.Name, 1.HouseNumber, 1.Birthdate
  - Entity.Id.Attribute: cust\_1234\_firstName
  - Hash the combination?
    - Hash("cust" + 1234 + "firstName") → f7941de89a66fbaa082db5d9b255a57d
    - Might have collisions between hashed values
    - No range searching
    - Improved distribution



# Key Design

- TIPS
  - Use meaningful and unambiguous naming components
    - 'cust' for customer, 'srep' for sales representative
  - Use range-based components if you need to retrieve ranges of values
    - Dates, Ids, etc.
  - Use a common delimiter when appending components
  - Keep keys as short as possible
- Well-designed keys minimize code required to access values
- Range key example for website visits
  - type\_date\_seqnum\_attribute
  - visit\_123116\_1\_sessionId, visit\_123116\_sessionId, etc.
- Make sure you account for the limitations of the KVDB you use
  - Byte limitations, data types
- Keep in mind your partitioning scheme when defining key structure

# Values

- Strong typing NOT required
  - “31 Spooner Street, Quahog, RI”
  - (“31 Spooner Street”, “Quahog”, “RI”)
  - { “street”:”31 Spooner Street”, “City”:”Quahog”, “State”:”RI” }
- TIP: Make implementation choices that lead to *some* restrictions
- Can’t search values, but can create a search index based on values
  - ‘IL’: (‘cust.1234.state’, ‘cust.2468.state’, ‘store.1011.state’)
  - Index the values and keys that have that value
  - Not necessarily efficient, but could improve scans
- Structured data can help reduce latency
  - Store commonly associated values together, like {firstName:Don, lastName: Sawyer}

# Data Samples (Again)

Key	Value
cust_1234_name	"Peter Griffin"
cust_1234_address	"31 Spooner St. Quahog, RI"
cust_1234_nameAddress	{"name": "Peter Griffin", "address": "31 Spooner St. Quahog, RI"}
cust_1234_profileImage	
cust_2468_name	"Homer Simpson"
cust_2468_address	"742 Evergreen Terrace, Springfield, ???"
cust_2468_nameAddress	{"name": "Homer Simpson", "address": "742 Evergreen Terrace, Springfield, ???"}
cust_2468_profileImage	

# Other Definitions

- Namespace
  - Collection of key-value pairs (aka set, collection, bucket)
  - Separate data in namespaces like customer, orders, products
- Partition
  - Organize data based on keys
    - Customer 1-10000, 10001-20000, ...
    - Different servers manage different partitions (multiple servers hold replicas of partitions for distributed reading)
    - You must understand your partitions and size in each
      - If customers are organized by name, there may be many more S's than Z's
    - Could use hash function to create hashed keys for better distribution
- Partition Key
  - The key used to determine which partition should hold a data value
- TTL (Time to Live)
  - Data goes away after a specific period of time

# Improving Performance

- Well-designed keys
- Structured values for fewer lookups
- Copies of data (denormalization) to assist usage
  - firstName, lastName, address, firstLastNameAddress
  - Combining data puts data in the same storage block (I/O reduction)
  - No need for joins
- If structured values are large, consider a document database

# Limitations

- Can only look up values by key
  - Some KVDBs have some version of search
  - Riak indexes values as well as keys
  - Use secondary indexes (if supported) or inverted indexes
- Range queries may not be supported
  - Ordered key-value db allows for this
- No comparable query language to SQL for RDBMS's
  - Cassandra has CQL
  - If storing JSON/XML, could integrate with Solr/Lucene for text search



# Key/Value: Useful When

- Unstructured data is required
- High performance read/writes
- Value is fully identifiable via key alone
- Value is not dependent on other values
- Values simplistic in structure or binary
- Simple query patterns (insert, select, delete only)
  - Ease of storage & retrieval more important than complex data structures
  - Update might be available
- Values are manipulated at application layer

# Key/Value: Not Useful When

- Need to search or filter data within the stored value
- Relationships exist between different key-value entries (joins)
- Multiple keys' values need to be updated in single transaction
- Multiple keys need to be modified in a single operation
- Schema consistency across values is required
- Partial updates are required (single attribute in a value)

# Document Databases

# Data Samples

## JSON Document

```
{
  "addresses": [
    {
      "city": "Quahog",
      "is_primary": "true",
      "number": "31",
      "state": "RI",
      "street": "Spooners St.",
      "type": "home"
    }
  ],
  "age": 42,
  "children": [
    {
      "name": "Meg Griffin",
      "type": "daughter"
    },
    {
      "name": "Brian Griffin",
      "type": "son"
    },
    {
      "name": "Stewie Griffin",
      "type": "son"
    }
  ],
  "first_name": "Peter",
  "last_name": "Griffin",
  "spouses": [
    {
      "name": "Lois Griffin"
    }
  ]
}
```

## XML Document

```
<Person>
  <Addresses>
    <Address IsPrimary="True" Type="home">
      <City>Quahog</City>
      <Number>31</Number>
      <State>RI</State>
      <Street>Spooners St.</Street>
    </Address>
  </Addresses>
  <Age>42</Age>
  <Children>
    <Child Type="daughter">
      <Name>Meg Griffin</Name>
    </Child>
    <Child Type="son">
      <Name>Brian Griffin</Name>
    </Child>
    <Child Type="son">
      <Name>Stewie Griffin</Name>
    </Child>
  </Children>
  <FirstName>Peter</FirstName>
  <LastName>Griffin</LastName>
  <Spouses>
    <Spouse>
      <Name>Lois Griffin</Name>
    </Spouse>
  </Spouses>
</Person>
```

# Design/Features

- Designed for scalability
- Provides flexibility about the structure of documents
  - Not all documents need to have the same structure
  - Can be completely unrelated documents (not advised)
    - E.G.: customer, sales, clickstream, logs
- Designed to accommodate variations and schema evolution
  - Hence, try to avoid explicit schema definitions like a RDBMS
- Schemaless
  - Schema specification not *required* (not schema on write)
  - Allows adding k-v pairs to documents
  - Application code must enforce rules about data
- For performance, think about balancing normalization vs. denormalization

# Document Collections

- Collection is a group/list of *related* documents
  - Documents don't have to have the same structure, but should share some common structure
- Tips
  - Avoid highly abstract entity types
    - Filtering collections is slower than working with multiple collections
    - Collections are stored near each other on disk, so you could end up reading a lot of data that needs to be filtered
  - You could index mixed entities, but won't necessarily be faster
    - If the entities are in separate collections, it might be faster to just scan the collection than read an index from disk
    - Indexing consumes resource to keep updated
  - Can use indexes on key terms instead scanning entire documents for an attribute value
  - Code for manipulating collections should apply to most/all documents
  - Use document subtypes when entities are aggregated/share substantial code
    - Products (appliances, clothing, music, toys)

# Normalization vs. Denormalization

- REMEMBER: DocDB is being used for its scalability
- Goal: keep data frequently used together in the document
  - Larger documents can lead to fewer documents retrieved in a block
- Consider queries your application will issue to the DB

## Joins Look Like (in application code):

```
for transactions in {Transaction collection query}
  for products in {Product collection query}:
    do something with transactions/products
```

Normalized Documents

Transactions Collection

```
{
  "transaction_id": 12345,
  "transaction_timestamp": 1420000000,
  "line_items": [
    { "upc": 1234, "name": "16.0L", "quantity": 1 },
    { "upc": 5678, "name": "8.0L", "quantity": 1 } ]
}
```

Products Collection

```
{
  "upc": 1234,
  "name": "16.0L",
  "category": "Beverages",
  "weight_kg": 1000,
},
{
  "upc": 5678,
  "name": "8.0L",
  "category": "Electronics",
  "weight_kg": 500
}
```

To get product names for a sale:

1. Look up transaction\_id [12345] in Transactions collection.
2. Get line\_items with upc values
3. Look up products by UPC [1234, 5678]
4. Loop through products and get product name

Looking up in two collections => extra disk I/O (think about scale here)

Denormalizing

Normalized

```
{
  "transaction_id": 12345,
  "transaction_timestamp": 1420000000,
  "line_items": [
    { "upc": 1234, "name": "16.0L", "quantity": 1 },
    { "upc": 5678, "name": "8.0L", "quantity": 1 } ]
}
```

Denormalized

```
{
  "transaction_id": 12345,
  "transaction_timestamp": 1420000000,
  "line_items": [
    {
      "upc": 1234,
      "name": "16.0L",
      "category": "Beverages",
      "weight_kg": 1000,
      "quantity": 1
    },
    {
      "upc": 5678,
      "name": "8.0L",
      "category": "Electronics",
      "weight_kg": 500,
      "quantity": 1
    }
  ]
}
```

# Operations on Document Databases

- Insert, update, delete, retrieve
- No standard data manipulation language (MongoDB Examples Below)
  - `db.product.insert({"product_id": 1234, "cost": 10.00})`
  - `db.product.remove({"product_id": 1234})`
  - `db.product.remove({"cost": {"$gte": 10.00}})`
  - `db.product.update({"product_id": 1234}, {$set {"cost": 9.99}})`
  - `db.product.update({"product_id": 1234}, {$set {"weight": 1.0}})`
  - `db.product.find({"product_id": 1234})`
- Tips
  - Include unique identifier with each document
  - Usually more efficient to bulk insert instead of [individual inserts](#)
  - Be careful when deleting documents with references to other documents



# Individual Inserts vs. Bulk Insert

## Individual Inserts

```
db.product.insert(  
    {"product_id": 1234,  
     "cost": 10.00} )  
  
db.product.insert(  
    {"product_id": 1235,  
     "cost": 1.50} )  
  
db.product.insert(  
    {"product_id": 1236,  
     "cost": 100.99} )
```

## Bulk Insert

```
db.product.insert( [  
    {"product_id": 1234,  
     "cost": 10.00},  
  
    {"product_id": 1235,  
     "cost": 1.50},  
  
    {"product_id": 1236,  
     "cost": 100.99} ] )
```

# Design: Partitioning

- Uses horizontal partitioning, not vertical
  - AKA Sharding
  - Divide database by documents (rows in RDBMS)
- Enables DB to scale horizontally
- Need to select a shard key + partitioning method
  - Key must exist in all documents
- Partitioning methods
  - Range (dates, numbers, alphabetic)
  - Hash (distribute keys evenly in partitions)
  - List (partition by lists of entities)
    - P1: Clothing | P2: Electronics, Toys, Office Supplies | P3: Grocery, Pharmacy

# More Design Considerations

- Physical model: planning for mutability
  - On creation a document is allocated space + room for growth
  - If the document grows > block, might be moved to another block
    - Creating a document w/ sufficient space will help avoid this
- Indexing
  - Too few => poor read performance
  - Too many => poor write performance
  - Read-heavy => index most/all fields
  - Write-heavy => focus on essential indexes (keys + identifiers of relations)
  - TIP: EXPERIMENT and iterate

# Modeling Relations

- Many-to-many
  - Use two collections
  - Each collection maintains a list of identifiers of related documents
  - Data integrity: be careful when updating many-to-many
- Hierarchies (parent-child, taxonomies)
  - Option: reference parent id of the parent
    - When you often need to traverse upwards
  - Option: reference child ids of children
    - When you often need to traverse downward
  - Option: list all ancestors (in path order)
    - When you need to know full path (single read to get full tree)

# Reference Slides

Slides referenced from earlier slides

# Normalized Documents

## Transactions Collection

```
{
  "transaction_id": 111111,
  "transaction_timestamp": 1481952866135,
  "line_items": [
    { "upc": 1234, "cost": 10.00, "quantity": 3 },
    { "upc": 2468, "cost": 9.99, "quantity": 1 }
  ]
}
```

### To get product names for a sale:

1. Look up transaction\_id (111111) in Transactions collection
  1. Get line\_items with upc values
2. Look up products by UPC [1234, 2468]
  1. Loop through products and get product name

*Looking up in two collections => extra disk I/O (think about scale here)*

## Products Collection

```
{
  "upc": 1234,
  "name": "Antitrust",
  "category": "Movies",
  "media_type": "DVD",
}

{
  "upc": 2468,
  "name": "3 ft micro USB cable",
  "category": "Electronics",
  "brand": "Belkin"
}
```

# Denormalizing

## Normalized

```
{
  "transaction_id": 111111,
  "transaction_timestamp": 1481952866135,
  "line_items": [
    { "upc": 1234, "cost": 10.00, "quantity": 3 },
    { "upc": 2468, "cost": 9.99, "quantity": 1 }
  ]
}

{
  "upc": 1234,
  "name": "Antitrust",
  "category": "Movies",
  "media_type": "DVD",
}

{
  "upc": 2468,
  "name": "3 ft micro USB cable",
  "category": "Electronics",
  "brand": "Belkin"
}
```

## Denormalized

```
{
  "transaction_id": 111111,
  "transaction_timestamp": 1481952866135,
  "line_items": [
    {
      "cost": 10.00,
      "quantity": 3,
      "product": {
        "upc": 1234,
        "name": "Antitrust",
        "category": "Movies",
        "media_type": "DVD"
      }
    },
    {
      "cost": 9.99,
      "quantity": 1,
      "product": {
        "upc": 2468,
        "name": "3 ft micro USB cable",
        "category": "Electronics",
        "brand": "Belkin"
      }
    }
  ]
}
```

# Use Cases

Some use cases where a document database could be used.



Give one reason why you would choose to move the following document data to a key-value database. Given an example of the schema design for the key-value record(s) for optimal performance.

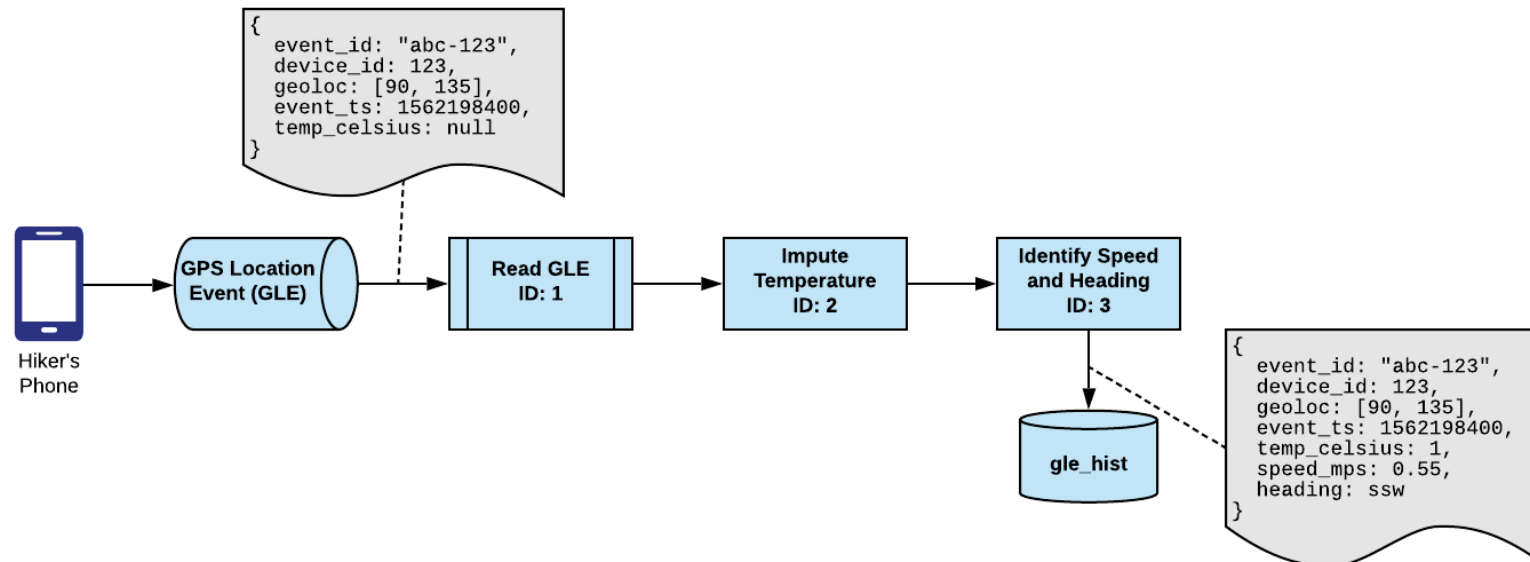
Document Collection: Hockey Scores

```
{  
  "game_id" : 1,  
  "home_score" : 4,  
  "visitor_score" : 3  
}
```

# Use Case #1: Metric Tracking

Design a database that can track provenance about your data pipeline. Use the example below.

## Hiking App Event Tracking Data Pipeline



# Check for Duplicates in a Realtime Stream

Realtime streams are often built for extremely fast processing of data coming at a high velocity. Many of the technologies use don't guarantee "exactly once" processing, rather "at least once" processing. How can you solve this potential duplication problem with K-V stores?